

SQL Notes

Asaad Cheema

SQL Commands and Operations

Below are detailed SQL commands based on the provided examples:

Stopping Database Alterations

To make a database read-only:

```
1 ALTER DATABASE myDB READ_ONLY = 1;
```

Creating a Table

Example of creating a table named `employees`:

```
1 CREATE TABLE employees (  
2     employee_id INT,  
3     first_name VARCHAR(50),  
4     last_name VARCHAR(50),  
5     hourly_pay DECIMAL(5,2),  
6     hire_date DATE  
7 );
```

Renaming a Table

Rename the table `employee` to `worker`:

```
1 RENAME TABLE employee TO worker;
```

Altering a Table

Adding a Column

Add a `phone_number` column to `employees`:

```
1 ALTER TABLE employees  
2 ADD phone_number VARCHAR(15);
```

Renaming a Column

Rename `phone_number` to `email`:

```
1 ALTER TABLE employees  
2 RENAME COLUMN phone_number TO email;
```

Modifying a Column

Change the datatype of email to VARCHAR(100):

```
1 ALTER TABLE employees
2 MODIFY COLUMN email VARCHAR(100);
```

Move the email column to be after last_name:

```
1 ALTER TABLE employees
2 MODIFY COLUMN email VARCHAR(100)
3 AFTER last_name;
```

Dropping a Column

Remove the email column:

```
1 ALTER TABLE employees
2 DROP COLUMN email;
```

Inserting Data

Inserting All Columns

Insert a complete record into employees:

```
1 INSERT INTO employees
2 VALUES (1, 'squid', 'krab', 5.50, '2023-02-01');
```

Inserting Specific Columns

Insert a record with specific columns:

```
1 INSERT INTO employees (employee_id, first_name, last_name, hire_date)
2 VALUES (1, 'squid', 'krab', '2023-02-01');
```

Selecting Data

Selecting All Records

Retrieve all records from employees:

```
1 SELECT * FROM employees;
```

Selecting Records with Conditions

Find records where hire_date is NULL:

```
1 SELECT * FROM employees
2 WHERE hire_date IS NULL;
```

Updating Data

Update the hourly_pay for an employee with employee_id = 1:

```
1 UPDATE employees
2 SET hourly_pay = 10.25
3 WHERE employee_id = 1;
```

Deleting Data

Delete a record with `employee_id = 6`:

```
1 DELETE FROM employees
2 WHERE employee_id = 6;
```

Transaction Control

Disabling Autocommit

Disable AUTOCOMMIT:

```
1 SET AUTOCOMMIT = OFF;
```

Committing Changes

Commit all changes made during the transaction:

```
1 COMMIT;
```

Rolling Back Changes

Roll back to the last committed state:

```
1 ROLLBACK;
```

Unique Constraints

Creating a Unique Column During Table Creation

A unique column ensures that no duplicate values exist in the column. Example:

```
1 CREATE TABLE products (
2     product_id INT,
3     product_name VARCHAR(100) UNIQUE,
4     price DECIMAL(10,2)
5 );
```

Adding a Unique Constraint to an Existing Table

To make an existing column unique after table creation:

```
1 ALTER TABLE products
2 ADD CONSTRAINT unique_product_name UNIQUE (product_name);
```

Example Usage

Consider a table `products` where we ensure that `product_name` is unique:

```
1 -- Add unique constraint to an existing column
2 ALTER TABLE products
3 ADD CONSTRAINT unique_product_name UNIQUE (product_name);
4
5 -- Attempting to insert duplicate product names will fail
6 INSERT INTO products (product_id, product_name, price)
7 VALUES (1, 'Laptop', 999.99);
8
9 -- This will throw an error as 'Laptop' already exists
10 INSERT INTO products (product_id, product_name, price)
11 VALUES (2, 'Laptop', 899.99);
```

NOT NULL Constraint

Defining a Column as NOT NULL During Table Creation

The NOT NULL constraint ensures that a column cannot have NULL values. Example:

```
1 CREATE TABLE orders (  
2     order_id INT NOT NULL,  
3     customer_name VARCHAR(100) NOT NULL,  
4     order_date DATE  
5 );
```

Adding NOT NULL Constraint to an Existing Column

To make an existing column NOT NULL:

```
1 ALTER TABLE orders  
2 MODIFY COLUMN customer_name VARCHAR(100) NOT NULL;
```

Example Usage

Consider a table `orders` where the `customer_name` column must always have a value:

```
1 -- Add a NOT NULL constraint  
2 ALTER TABLE orders  
3 MODIFY COLUMN customer_name VARCHAR(100) NOT NULL;  
4  
5 -- Attempting to insert a row without customer_name will fail  
6 INSERT INTO orders (order_id, order_date)  
7 VALUES (1, '2023-01-01'); -- Error: customer_name cannot be NULL
```

CHECK Constraint

Defining a Column with a CHECK Constraint During Table Creation

The CHECK constraint ensures that a column's values meet a specific condition. Example:

```
1 CREATE TABLE employees (  
2     employee_id INT,  
3     employee_name VARCHAR(100),  
4     age INT CHECK (age >= 18),  
5     salary DECIMAL(10,2) CHECK (salary > 0)  
6 );
```

Adding a CHECK Constraint to an Existing Column

To add a CHECK constraint to an existing column:

```
1 ALTER TABLE employees  
2 ADD CONSTRAINT check_age CHECK (age >= 18);
```

Example Usage

Consider a table `employees` where the `age` column must always have a value of 18 or greater:

```
1 -- Add a CHECK constraint  
2 ALTER TABLE employees  
3 ADD CONSTRAINT check_age CHECK (age >= 18);  
4  
5 -- Attempting to insert a row with age less than 18 will fail  
6 INSERT INTO employees (employee_id, employee_name, age, salary)  
7 VALUES (1, 'John Doe', 17, 3000.00); -- Error: CHECK constraint failed
```

DEFAULT Constraint

Defining a Column with a DEFAULT Value During Table Creation

The DEFAULT constraint assigns a default value to a column if no value is specified during insertion. Example:

```
1 CREATE TABLE employees (  
2     employee_id INT,  
3     employee_name VARCHAR(100),  
4     department VARCHAR(50) DEFAULT 'General',  
5     salary DECIMAL(10,2) DEFAULT 3000.00  
6 );
```

Adding a DEFAULT Constraint to an Existing Column

To add a DEFAULT constraint to an existing column:

```
1 ALTER TABLE employees  
2 ALTER COLUMN department SET DEFAULT 'General';
```

Example Usage

Consider a table `employees` where the `department` column has a default value of 'General':

```
1 -- Insert without specifying the department  
2 INSERT INTO employees (employee_id, employee_name, salary)  
3 VALUES (1, 'John Doe', 5000.00);  
4  
5 -- The department will automatically be set to 'General'.  
6 SELECT * FROM employees;
```

PRIMARY KEY Constraint

Defining a Primary Key During Table Creation

The PRIMARY KEY constraint uniquely identifies each record in a table. Example:

```
1 CREATE TABLE orders (  
2     order_id INT PRIMARY KEY,  
3     customer_name VARCHAR(100),  
4     order_date DATE  
5 );
```

Adding a Primary Key to an Existing Table

To add a PRIMARY KEY constraint to an existing table:

```
1 ALTER TABLE orders  
2 ADD CONSTRAINT pk_order_id PRIMARY KEY (order_id);
```

Example Usage

Consider a table `orders` where the `order_id` column is the primary key:

```
1 -- Insert valid rows  
2 INSERT INTO orders (order_id, customer_name, order_date)  
3 VALUES (1, 'Alice', '2023-01-01');  
4  
5 -- Attempting to insert a duplicate primary key will fail  
6 INSERT INTO orders (order_id, customer_name, order_date)  
7 VALUES (1, 'Bob', '2023-01-02'); -- Error: Duplicate entry for primary key
```

FOREIGN KEY Constraints

Defining a FOREIGN KEY During Table Creation

The FOREIGN KEY constraint establishes a relationship between two tables. Example:

```
1 CREATE TABLE orders (  
2     order_id INT,  
3     customer_id INT,  
4     order_date DATE,  
5     PRIMARY KEY (order_id),  
6     FOREIGN KEY (customer_id) REFERENCES customers(customer_id)  
7 );
```

Adding a FOREIGN KEY to an Existing Table

To add a FOREIGN KEY constraint to an existing table:

```
1 ALTER TABLE orders  
2 ADD CONSTRAINT fk_customer FOREIGN KEY (customer_id)  
3 REFERENCES customers(customer_id);
```

Example Usage

Consider a table `orders` linked to a table `customers` through a foreign key:

```
1 -- Insert records into customers  
2 INSERT INTO customers (customer_id, customer_name)  
3 VALUES (1, 'John Doe'), (2, 'Jane Smith');  
4  
5 -- Insert records into orders  
6 INSERT INTO orders (order_id, customer_id, order_date)  
7 VALUES (101, 1, '2023-01-01'), (102, 2, '2023-01-02');
```

SQL Joins

INNER JOIN

Retrieve records with matching values in both tables:

```
1 SELECT orders.order_id, customers.customer_name  
2 FROM orders  
3 INNER JOIN customers ON orders.customer_id = customers.customer_id;
```

LEFT JOIN

Retrieve all records from the left table and matching records from the right table:

```
1 SELECT orders.order_id, customers.customer_name  
2 FROM orders  
3 LEFT JOIN customers ON orders.customer_id = customers.customer_id;
```

RIGHT JOIN

Retrieve all records from the right table and matching records from the left table:

```
1 SELECT orders.order_id, customers.customer_name  
2 FROM orders  
3 RIGHT JOIN customers ON orders.customer_id = customers.customer_id;
```

FULL JOIN

Retrieve all records when there is a match in either table:

```
1 SELECT orders.order_id, customers.customer_name
2 FROM orders
3 FULL OUTER JOIN customers ON orders.customer_id = customers.customer_id;
```

Logical Operators: AND, OR, NOT

Logical operators are used in SQL to combine conditions in the **WHERE** clause.

AND

The AND operator combines conditions and returns rows where all conditions are true.

```
1 SELECT * FROM employees
2 WHERE department = 'HR' AND employee_name = 'Alice';
```

OR

The OR operator combines conditions and returns rows where at least one condition is true.

```
1 SELECT * FROM employees
2 WHERE department = 'HR' OR department = 'IT';
```

NOT

The NOT operator negates a condition and returns rows where the condition is false.

```
1 SELECT * FROM employees
2 WHERE NOT department = 'HR';
```

Combined Example

You can combine AND, OR, and NOT operators for complex conditions.

```
1 SELECT * FROM employees
2 WHERE (department = 'HR' OR department = 'IT')
3 AND NOT employee_name = 'Alice';
```

This query retrieves all employees from the HR or IT departments except those named Alice.

SQL Functions

SQL provides various functions to perform calculations on data. Some commonly used functions are:

Aggregate Functions

- **SUM()**: Calculates the total of a numeric column.
- **COUNT()**: Counts the number of rows.
- **AVG()**: Calculates the average value of a numeric column.
- **MAX()**: Returns the maximum value in a column.
- **MIN()**: Returns the minimum value in a column.

Examples

```
1 -- Calculate the total salary of all employees
2 SELECT SUM(salary) AS total_salary FROM employees;
3
4 -- Count the number of employees in each department
5 SELECT department, COUNT(*) AS num_employees
6 FROM employees
7 GROUP BY department;
8
9 -- Find the average salary in the company
10 SELECT AVG(salary) AS average_salary FROM employees;
11
12 -- Get the maximum and minimum salaries in the company
13 SELECT MAX(salary) AS highest_salary, MIN(salary) AS lowest_salary
14 FROM employees;
```

Wildcards with LIKE

Wildcards are used with the LIKE operator to search for patterns in text columns.

Wildcards

- %: Represents zero or more characters.
- _: Represents a single character.

Examples

```
1 -- Find employees whose names start with 'A'
2 SELECT * FROM employees
3 WHERE employee_name LIKE 'A%';
4
5 -- Find employees whose names end with 'n'
6 SELECT * FROM employees
7 WHERE employee_name LIKE '%n';
8
9 -- Find employees whose names contain 'ali'
10 SELECT * FROM employees
11 WHERE employee_name LIKE '%ali%';
12
13 -- Find employees whose names have 'a' as the second character
14 SELECT * FROM employees
15 WHERE employee_name LIKE '_a%';
```

Combined Example

Using functions and wildcards together:

```
1 -- Find the total salary of employees whose names start with 'A'
2 SELECT SUM(salary) AS total_salary
3 FROM employees
4 WHERE employee_name LIKE 'A%';
5
6 -- Count the number of employees in departments starting with 'S'
7 SELECT COUNT(*) AS num_employees
8 FROM employees
9 WHERE department LIKE 'S%';
```

LIMIT Clause

The LIMIT clause is used to restrict the number of rows returned by a query. It can be combined with the OFFSET keyword to skip a specific number of rows before starting to return rows.

Examples

```
1 -- Retrieve the first 5 rows from the employees table
2 SELECT * FROM employees
3 LIMIT 5;
4
5 -- Retrieve the next 5 rows after skipping the first 5 rows
6 SELECT * FROM employees
7 LIMIT 5 OFFSET 5;
8
9 -- Retrieve rows 11 to 20 from the employees table
10 SELECT * FROM employees
11 LIMIT 10 OFFSET 10;
12
13 -- MySQL-specific syntax to retrieve rows 11 to 20
14 SELECT * FROM employees
15 LIMIT 10, 10;
```

Note

In MySQL, `LIMIT 10 OFFSET 10` is equivalent to `LIMIT 10, 10`. While the former is widely supported across various databases, the latter is specific to MySQL.

ORDER BY Clause

The `ORDER BY` clause is used to sort the result set in ascending or descending order.

Examples

```
1 -- Retrieve all employees sorted by their names in ascending order
2 SELECT * FROM employees
3 ORDER BY employee_name ASC;
4
5 -- Retrieve employees sorted by salary in descending order
6 SELECT * FROM employees
7 ORDER BY salary DESC;
```

GROUP BY Clause

The `GROUP BY` clause groups rows that have the same values in specified columns. It is often used with aggregate functions (Where is replaced by Having).

Examples

```
1 -- Count the number of employees in each department
2 SELECT department, COUNT(*) AS num_employees
3 FROM employees
4 GROUP BY department;
5
6 -- Calculate the total salary for each department
7 SELECT department, SUM(salary) AS total_salary
8 FROM employees
9 GROUP BY department;
```

Indexes in MySQL

Indexes are used to optimize the performance of queries by allowing faster data retrieval. MySQL supports several types of indexes:

Types of Indexes

- **PRIMARY KEY:** A unique index where each value must be unique and cannot be NULL.
- **UNIQUE:** Ensures all values in the indexed column are unique.
- **FULLTEXT:** Used for text searching.
- **INDEX (Non-unique):** Speeds up data retrieval but does not enforce uniqueness.

Creating Indexes

```
1 -- Create a simple index
2 CREATE INDEX idx_employee_name ON employees(employee_name);
3
4 -- Create a unique index
5 CREATE UNIQUE INDEX idx_unique_department ON departments(department_name);
6
7 -- Add an index when creating a table
8 CREATE TABLE employees (
9     employee_id INT PRIMARY KEY,
10    employee_name VARCHAR(100),
11    department_id INT,
12    INDEX idx_department_id (department_id)
13 );
```

Dropping Indexes

```
1 -- Drop an index
2 DROP INDEX idx_employee_name ON employees;
```

Subqueries in MySQL

Subqueries are nested queries used within another SQL query to perform operations that depend on the results of the nested query.

Types of Subqueries

- **Single-row Subqueries:** Return a single value.
- **Multi-row Subqueries:** Return multiple rows.
- **Correlated Subqueries:** Refer to columns from the outer query.

Examples

```
1 -- Single-row subquery to find the employee with the highest salary
2 SELECT * FROM employees
3 WHERE salary = (SELECT MAX(salary) FROM employees);
4
5 -- Multi-row subquery to find employees in specific departments
6 SELECT employee_name FROM employees
7 WHERE department_id IN (SELECT department_id FROM departments WHERE location = 'New York');
8
9 -- Correlated subquery to find employees earning above the average salary in their
10 department
11 SELECT employee_name, salary FROM employees e1
12 WHERE salary > (
13     SELECT AVG(salary) FROM employees e2
14     WHERE e1.department_id = e2.department_id
15 );
```

Views in MySQL

Views are virtual tables based on the result set of a query. They do not store data but simplify complex queries and enhance security by controlling data access.

Creating Views

```
1 -- Create a view for employee details with department names
2 CREATE VIEW employee_department AS
3 SELECT e.employee_id, e.employee_name, d.department_name
4 FROM employees e
5 JOIN departments d
6 ON e.department_id = d.department_id;
```

Using Views

```
1 -- Query the view like a regular table
2 SELECT * FROM employee_department;
3
4 -- Use the view with filtering
5 SELECT employee_name, department_name
6 FROM employee_department
7 WHERE department_name = 'HR';
```

Dropping Views

```
1 -- Drop a view
2 DROP VIEW employee_department;
```

ROLLUP in MySQL

The ROLLUP operator is used in conjunction with the GROUP BY clause to generate subtotals and grand totals.

Example

```
1 -- Calculate total salary by department with a grand total
2 SELECT department_id, SUM(salary) AS total_salary
3 FROM employees
4 GROUP BY department_id WITH ROLLUP;
5
6 -- Output example:
7 -- | department_id | total_salary |
8 -- | 1             | 50000       |
9 -- | 2             | 75000       |
10 -- | NULL          | 125000      | -- Grand total
```

Notes

- The NULL value in the ROLLUP result represents the grand total.
- Useful for generating reports with aggregated data.

Triggers in MySQL

Triggers are database objects that automatically execute a specified SQL statement when a specific event occurs on a table.

Creating Triggers

```
1 -- Create a trigger to log changes to a table
2 CREATE TRIGGER after_employee_update
3 AFTER UPDATE ON employees
4 FOR EACH ROW
5 BEGIN
6     INSERT INTO employee_changes (employee_id, change_time, old_salary, new_salary)
7     VALUES (OLD.employee_id, NOW(), OLD.salary, NEW.salary);
8 END;
```

Using Triggers

```
1 -- Example: Update an employee's salary
2 UPDATE employees
3 SET salary = 60000
4 WHERE employee_id = 1;
5
6 -- The trigger logs the change in the employee_changes table.
7 SELECT * FROM employee_changes;
```

Dropping Triggers

```
1 -- Drop a trigger
2 DROP TRIGGER after_employee_update;
```

Notes

- Triggers can be defined for INSERT, UPDATE, or DELETE events.
- Each trigger is associated with a specific table.

Notes

- Triggers can be defined for INSERT, UPDATE, or DELETE events.
- Each trigger is associated with a specific table.

ON DELETE Clause in MySQL

The ON DELETE clause is used to define actions that should occur when a referenced row in a parent table is deleted. Common options include SET NULL and CASCADE.

SET NULL

If ON DELETE SET NULL is used, the foreign key column in the child table will be set to NULL when the referenced row in the parent table is deleted.

```
1 -- Example: Create tables with ON DELETE SET NULL
2 CREATE TABLE departments (
3     department_id INT PRIMARY KEY,
4     department_name VARCHAR(100)
5 );
6
7 CREATE TABLE employees (
8     employee_id INT PRIMARY KEY,
9     employee_name VARCHAR(100),
10    department_id INT,
11    FOREIGN KEY (department_id) REFERENCES departments(department_id)
12    ON DELETE SET NULL
13 );
14
15 -- Delete a department
```

```

16 DELETE FROM departments WHERE department_id = 1;
17
18 -- The department_id column in the employees table will be set to NULL.
19 SELECT * FROM employees;

```

CASCADE

If ON DELETE CASCADE is used, deleting a row in the parent table will automatically delete all related rows in the child table.

```

1 -- Example: Create tables with ON DELETE CASCADE
2 CREATE TABLE departments (
3     department_id INT PRIMARY KEY,
4     department_name VARCHAR(100)
5 );
6
7 CREATE TABLE employees (
8     employee_id INT PRIMARY KEY,
9     employee_name VARCHAR(100),
10    department_id INT,
11    FOREIGN KEY (department_id) REFERENCES departments(department_id)
12    ON DELETE CASCADE
13 );
14
15 -- Delete a department
16 DELETE FROM departments WHERE department_id = 2;
17
18 -- All employees in department_id = 2 will also be deleted.
19 SELECT * FROM employees;

```

Stored Procedures in MySQL

Stored procedures are precompiled SQL statements that can be executed as a single unit. They allow for complex operations and can accept parameters.

Creating a Stored Procedure

```

1 -- Create a stored procedure to calculate the total salary by department
2 CREATE PROCEDURE GetTotalSalaryByDepartment (IN dep_id INT, OUT total_salary DECIMAL
3     (10,2))
4 BEGIN
5     SELECT SUM(salary) INTO total_salary
6     FROM employees
7     WHERE department_id = dep_id;
8 END;

```

Using a Stored Procedure

```

1 -- Declare a variable to store the output
2 SET @total = 0;
3
4 -- Call the procedure
5 CALL GetTotalSalaryByDepartment(1, @total);
6
7 -- Display the result
8 SELECT @total AS TotalSalary;

```

Dropping a Stored Procedure

```

1 -- Drop a stored procedure
2 DROP PROCEDURE GetTotalSalaryByDepartment;

```

Notes

- Stored procedures can have input (**IN**), output (**OUT**), or input-output (**INOUT**) parameters.
- They are useful for encapsulating complex logic and reusing it across multiple queries.